

Foundations of Computer Science in Python

Lecture 8: NumPy & SciPy



Today's Plan

- Introduction to *NumPy* & *SciPy*
- Plotting
- Data Analysis



NumPy and SciPy

- **NumPy** and **SciPy** are packages for scientific computing that provide fast pre-compiled mathematical and numerical functions.
- The **NumPy** (Numeric Python) package provides basic routines for manipulating large arrays and matrices of numeric data.
- The **SciPy** (Scientific Python) package extends the functionality of NumPy with a large collection of useful algorithms, like minimization, regression, and other applied mathematical techniques

NumPy and SciPy

- **NumPy** and **SciPy** are open-source, and therefore provide a free Matlab alternative.
- The packages are popular among scientists, researchers and engineers who want to apply various mathematical methods on large datasets.

Read more here: <http://docs.scipy.org/doc/numpy/reference/>

Importing the required modules

- The packages we need are already included within the installation of Anaconda.
- Make sure you import any needed package at the beginning of your program in order to use its classes....
- For example:

```
import numpy as np, matplotlib, scipy
```



Alias

NumPy's main object is the **Array**

- NumPy's main object is the homogeneous multidimensional **array**.
- It is a table of elements (usually numbers), all of the **same type**, indexed by a tuple of positive integers.
- In Numpy dimensions are called **axes**.
- The number of axes is **rank**.
 - A **vector** is an array of rank 1
 - A 2D **matrix** is an array of rank 2.

The Array object - Creation

```
>>> import numpy as np
```

```
>>> a = np.array([0, 1, 2])
```

```
array([0,1,2])
```

1 x 3 array

0	1	2
---	---	---

```
>>> b = np.array([[0, 1, 2], [3, 4, 5]])
```

```
array([[0,1,2],  
       [3,4,5]])
```

2 x 3 array

0	1	2
3	4	5

Array – Attributes & Methods

```
>>> b = np.array([[0, 1, 2], [3, 4, 5]]) # Creates a 2 x 3 array
```

```
array([[0,1,2],  
       [3,4,5]])
```

```
>>> b.ndim #number of dimensions
```

```
2
```

```
>>> b.shape #dimension sizes
```

```
(2,3)
```

```
>>> len(b) # returns the size of the first dimension
```

```
2
```

```
>>> b.T      #Transpose
```

```
array([[0, 3],  
       [1, 4],  
       [2, 5]])
```

0	1	2
3	4	5

0	3
1	4
2	5

Creating an Array of a specific type

```
>>> a = np.array(range(10), float)
>>> a
array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9.])
>>> a.dtype
dtype('float64')
```

Creating an array of zeros or ones

```
>>> np.zeros(7, dtype=int)  
array([0, 0, 0, 0, 0, 0, 0])
```

```
>>> np.ones((2,3), dtype=float)  
array([[ 1.,  1.,  1.],  
       [ 1.,  1.,  1.]])
```

Creating arrays containing number sequences

```
>>> np.arange(10)  
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
>>> np.arange(2,7)  
array([2, 3, 4, 5, 6])
```

```
>>> np.arange(2,7,2) # start, end (exclusive), step  
array([2, 4, 6])
```

Like range in Python, but
returns a numpy array object

Creating arrays containing number sequences

```
>>> print (np.arange(1, 5.5, 0.5)) # start, end (exclusive), step  
[ 1.  1.5  2.  2.5  3.  3.5  4.  4.5  5.]
```

*# **slicing** is similar to standard lists*

```
>>> print (np.arange(1, 5.5, 0.5)[3:1:-1])  
[ 2.5  2.]
```

```
>>> print(np.linspace(1,5,9)) # start, end, number of points  
[ 1.  1.5  2.  2.5  3.  3.5  4.  4.5  5.]
```

Arrays with random values between 0 and 1

```
>>> print (np.random.rand(5))  
[ 0.53844313  0.3384871  0.5763536  0.29159273  0.43938366]
```

Indexing and slicing: similar to lists

```
>>> a = np.diag(np.arange(3))
array([[0, 0, 0],
       [0, 1, 0],
       [0, 0, 2]])
>>> a[1, 1]
1
>>> a[2, 1] = 0
>>> a[1, ] # a[1] also works
array([ 0, 1, 0])
>>> a[:,1]
array([ 0, 1, 0])
>>> a[2][1:]
array([0, 2])
```

**Slicing: specify which
rows/columns
to take by
A[rows, columns]**

Indexing and slicing: matrices

```
>>> a[0,3:5]  
array([3, 4])
```

```
>>> a[4:,4:]  
array([[44, 45],  
       [54, 55]])
```

```
>>> a[:,2]  
array([2, 12, 22, 32, 42, 52])
```

```
>>> a[2::2,::2]  
array([[20, 22, 24],  
       [40, 42, 44]])
```

0	1	2	3	4	5
10	11	12	13	14	15
20	21	22	23	24	25
30	31	32	33	34	35
40	41	42	43	44	45
50	51	52	53	54	55

Reshaping an Array

```
>>> a
```

```
array([ 0., 1., 2., 3., 4., 5., 6., 7., 8., 9.])
```

0.	1.	2.	3.	4.	5.	6.	7.	8.	9.
----	----	----	----	----	----	----	----	----	----

```
>>> a = a.reshape((5, 2))
```

```
>>> a
```

```
array([[ 0.,  1.],  
       [ 2.,  3.],  
       [ 4.,  5.],  
       [ 6.,  7.],  
       [ 8.,  9.]])
```

```
>>> a.shape
```

```
(5, 2)
```



Reshape

0.	1.
2.	3.
4.	5.
6.	7.
8.	9.

Example

```
>>> x = np.arange(10)
>>> x[2]
2
>>> x[-2]
8
>>> x.shape = (2,5) # Setting x's dimensions to (2,5)
>>> x[1,3]
8
>>> x[1,-1]
9
>>> x[0]
array([0, 1, 2, 3, 4])
```


Array Arithmetic

Arithmetic operators on arrays apply *elementwise*. A new array is created and filled with the result.

```
>>> x = np.array([1,5,2])
```

```
>>> y = np.array([7,4,1])
```

```
>>> x + y
```

```
array([8, 9, 3])
```

```
>>> x * y # element by element multiplication ! Use np.dot(x,y) for matrix multiplication.
```

```
array([ 7, 20,  2])
```

```
>>> x - y
```

```
array([-6,  1,  1])
```

```
>>> x / y
```

```
array([0, 1, 2])
```

```
>>> x % y
```

```
array([1, 1, 0])
```

Note that here, x and y have the same size


Broadcasting

- NumPy operations are usually done element-by-element which requires two arrays to have exactly the same shape:

```
>>> a = array([1.0,2.0,3.0])  
>>> b = array([2.0,2.0,2.0])  
>>> a * b  
array([ 2.,  4.,  6.])
```

- But this will also work:


```
>>> a = array([1.0,2.0,3.0])  
>>> b = 2.0  
>>> a * b  
array([ 2.,  4.,  6.])
```



Broadcasting

- Broadcasting is supported when the size of the *trailing* axes for both arrays in an operation is either the same or when one of them is one.

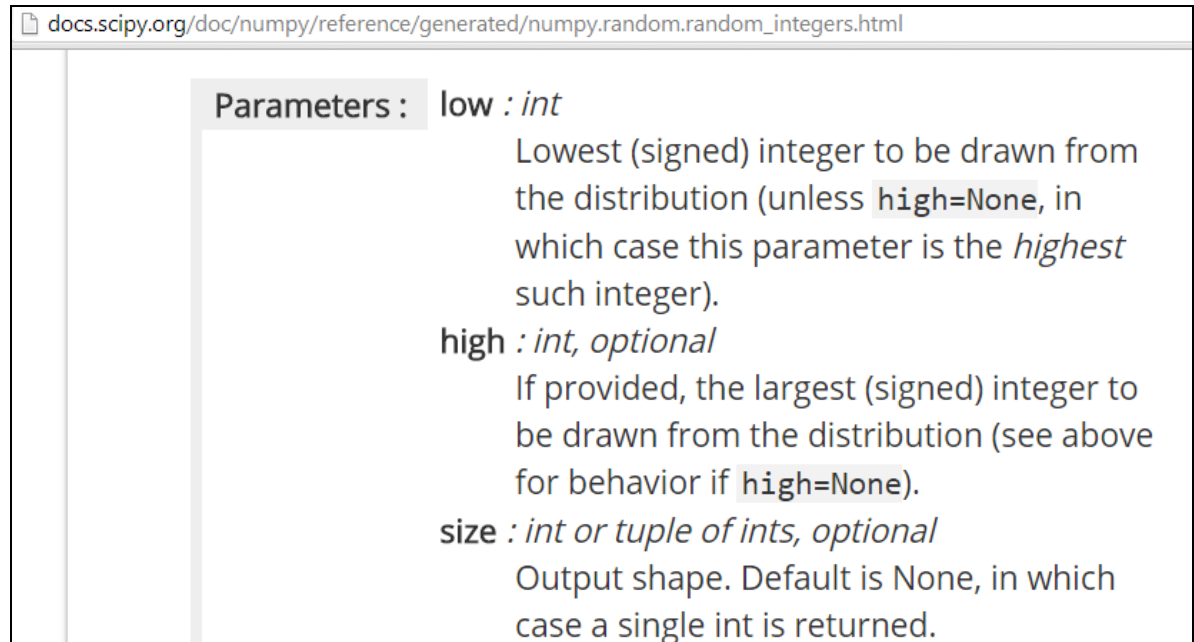
```
>>> a = array([[ 0.0,  0.0,  0.0],
...           [10.0,10.0,10.0],
...           [20.0,20.0,20.0],
...           [30.0,30.0,30.0]])
>>> b = array([1.0,2.0,3.0])
>>> a + b
array([[ 1.,  2.,  3.],
       [11., 12., 13.],
       [21., 22., 23.],
       [31., 32., 33.]])
```



Comparisons and Boolean operations

```
>>> a = np.random.random_integers(0, 10, 15)  
array([ 4,  9,  1,  7,  0,  9,  8, 10,  1,  0,  7,  7,  9,  5,  1])
```

Documentation - search
np.random.integers in
google



The screenshot shows a web browser window with the URL `docs.scipy.org/doc/numpy/reference/generated/numpy.random.random_integers.html`. The page content displays the parameters for the `random_integers` function. A 'Parameters:' label is followed by three entries: `low`, `high`, and `size`, each with its type and a descriptive text block.

Parameters :

- low** : *int*
Lowest (signed) integer to be drawn from the distribution (unless `high=None`, in which case this parameter is the *highest* such integer).
- high** : *int, optional*
If provided, the largest (signed) integer to be drawn from the distribution (see above for behavior if `high=None`).
- size** : *int or tuple of ints, optional*
Output shape. Default is `None`, in which case a single `int` is returned.

Comparisons and Boolean operations

```
>>> a = np.random.random_integers(0, 10, 15)
>>> b = np.random.random_integers(0, 10, 15)
>>> comp1 = a==b
>>> comp1
array([False, False, False, False, False,  True, False, False,
       False,  False, False, False, False, False], dtype=bool)
>>> comp1.any()
True
>>> comp1.all()
False
>>> comp1.nonzero()
(array([5], dtype=int64),)
```

Any is true?

Are all true?

Get the True indices

Fancy indexing: Logical indexing

```
>>> a = np.random.random_integers(0, 20, 15)  
array([10, 3, 8, 0, 19, 10, 11, 9, 10, 6, 0, 20, 12, 7, 14])  
>>> (a % 3 == 0)
```

What will be the result?
What will be the type of the new object?

```
array([False, True, False, True, False, False, False, True,  
False, True, True, False, True, False, False], dtype=bool)
```

Fancy indexing: Logical indexing

```
>>> a = np.random.random_integers(0, 20, 15)
array([10, 3, 8, 0, 19, 10, 11, 9, 10, 6, 0, 20, 12, 7, 14])
>>> (a % 3 == 0)
array([False,  True,  False,  True,  False,  False,  False,  True,
        False,  True,  True,  False,  True,  False,  False], dtype=bool)
>>> mask = (a % 3 == 0)
# extract sub-array: this is a copy!
>>> extract_from_a = a[mask]
array([ 3, 0, 9, 6, 0, 12])
# change sub-array
>>> a[mask] = -1
array([10, -1, 8, -1, 19, 10, 11, -1, 10, -1, -1, 20, -1, 7, 14])
```

Fancy indexing II

Indexing with array/list of integers

```
>>> a = np.arange(0, 100, 10)
```

note: [2, 3, 2, 4, 2] is a Python list

```
>>> a[[2, 3, 2, 4, 2]]
```


```
array([20, 30, 20, 40, 20])
```


Reductions

```
>>> x = np.array([[1, 1], [2, 2]])  
array([[1, 1],  
       [2, 2]])  
>>> x.sum() # works on the entire matrix  
6
```

```
>>> x.sum(axis=0) # sum over the columns (first dimension)  
array([3, 3])  
>>> x[:, 0].sum(), x[:, 1].sum()  
(3, 3)
```

```
>>> x.sum(axis=1) # sum over the rows (second dimension)  
array([2, 4])  
>>> x[0, :].sum(), x[1, :].sum()  
(2, 4)
```



**Also works with
x.min, x.max,
x.mean etc.**

Sorting along an axis

```
>>> a = np.array([[4, 3, 5], [1, 2, 1]])
```

```
>>> b = np.sort(a, axis=1)
```

```
>>> b
```

```
array([[3, 4, 5],  
       [1, 1, 2]])
```

```
>>> a.sort(axis=1)
```

```
>>> a
```



```
array([[3, 4, 5],  
       [1, 1, 2]])
```

Array stacking

- Enables concatenating two arrays vertically (***vstack***) or horizontally (***hstack***)

```
>>> a = np.array((1,2,3))
>>> b = np.array((2,3,4))
>>> np.hstack((a,b))
array([1, 2, 3, 2, 3, 4])
>>> a = np.array([[1],[2],[3]])
>>> b = np.array([[2],[3],[4]])
>>> np.hstack((a,b))
array([[1, 2],
       [2, 3],
       [3, 4]])
```



```
>>> a = np.array([1, 2, 3])
>>> b = np.array([2, 3, 4])
>>> np.vstack((a,b))
array([[1, 2, 3],
       [2, 3, 4]])
>>> a = np.array([[1], [2], [3]])
>>> b = np.array([[2], [3], [4]])
>>> np.vstack((a,b))
array([[1],
       [2],
       [3],
       [2],
       [3],
       [4]])
```



Plotting


matplotlib

"A plot is worth a thousand words"






- **matplotlib** is a python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive.
- You can generate plots, histograms, power spectra, bar charts, errorcharts, scatterplots, etc, with just a few lines of code.
- Check this out: <http://matplotlib.org/gallery.html>



Example 1

```
import numpy as np
import matplotlib.pyplot as plt

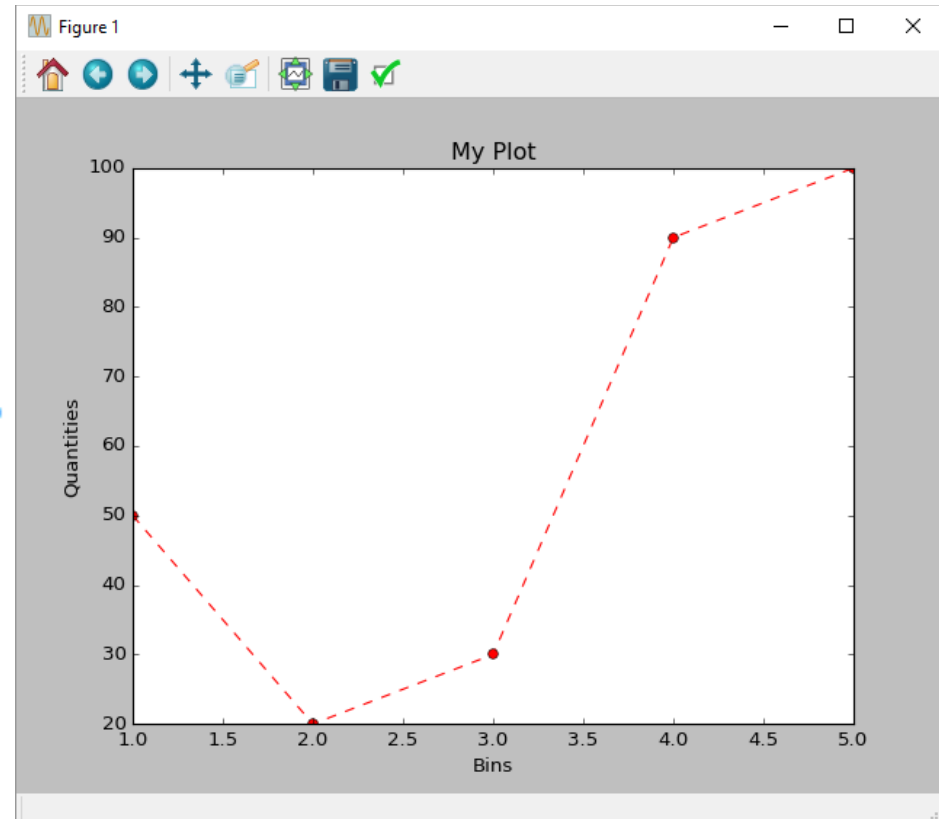
x = np.array[5 ,4 ,3 ,2 ,1)
y = np.array([100 ,90 ,30 ,20 ,50
plt.plot( x,y)
plt.show()
```



Example 1b

```
import numpy as np
import matplotlib.pyplot as plt

x = np.array ([1, 2, 3, 4, 5])
y = np.array ([50, 20, 30, 90, 100])
plt.plot (x,y,color='red',ls='--',marker='o')
plt.xlabel ('Bins')
plt.ylabel ('Quantities')
plt.title ('My Plot')
plt.show()
```




Example 2

```
import matplotlib.pyplot as plt  
import numpy as np
```

```
x = np.linspace(-np.pi, np.pi)  
y, z = np.cos(x), np.sin(x)
```

```
plt.plot(x, y, color = 'green')  
plt.plot(x, z, color = 'blue')  
plt.show()
```

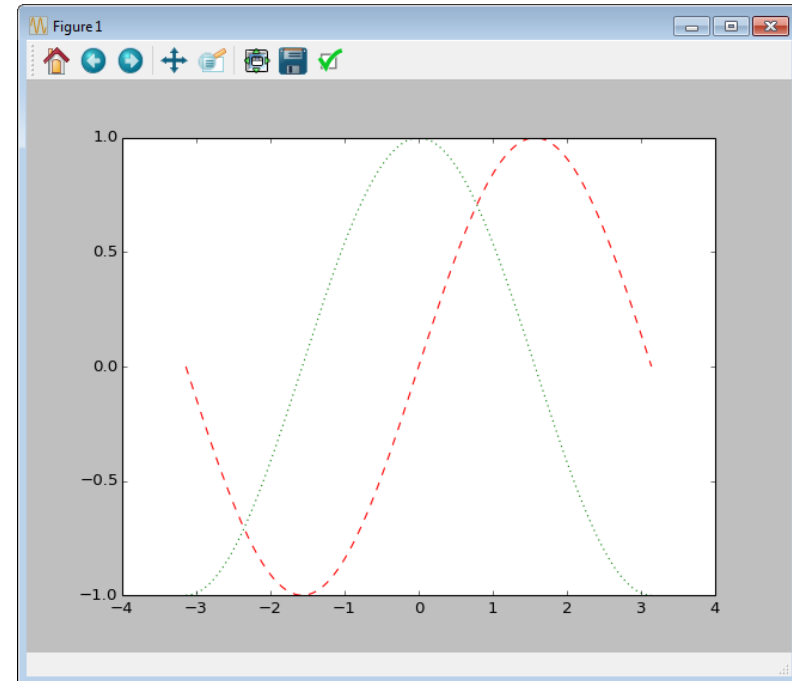


Example 2b

```
x = np.linspace(-np.pi, np.pi)
y, z = np.cos(x), np.sin(x)
plt.plot(x, y, color = 'green')
plt.plot(x, z, color = 'r--')
plt.show()
```

character	color
'b'	blue
'g'	green
'r'	red
'c'	cyan
'm'	magenta
'y'	yellow
'k'	black
'w'	white

character	description
'-'	solid line style
'--'	dashed line style
'-.'	dash-dot line style
':'	dotted line style
'.'	point marker
','	pixel marker
'o'	circle marker
'v'	triangle_down marker
'^'	triangle_up marker
'<'	triangle_left marker
'>'	triangle_right marker
'1'	tri_down marker
'2'	tri_up marker
'3'	tri_left marker
'4'	tri_right marker



Data analysis using Python

Real data analysis example

- Look at the file inflammation-01.csv (download the lecture's code)
- Tables as CSV files
 - Text files
 - Each row holds the same number of columns
 - Values are separated by commas

```
0,0,1,3,1,2,4,7,8,3,3,3,10,5,7,4,7,7,12,18,6,13,11,11,7,7,4,6,8,8,4,4,5,7,3,4,2,3,0,0
0,1,2,1,2,1,3,2,2,6,10,11,5,9,4,4,7,16,8,6,18,4,12,5,12,7,11,5,11,3,3,5,4,4,5,5,1,1,0,1
0,1,1,3,3,2,6,2,5,9,5,7,4,5,4,15,5,11,9,10,19,14,12,17,7,12,11,7,4,2,10,5,4,2,2,3,2,2,1,1
0,0,2,0,4,2,2,1,6,7,10,7,9,13,8,8,15,10,10,7,17,4,4,7,6,15,6,4,9,11,3,5,6,3,3,4,2,3,2,1
0,1,1,3,3,1,3,5,2,4,4,7,6,5,3,10,8,10,6,17,9,14,9,7,13,9,12,6,7,7,9,6,3,2,2,4,2,0,1,1
```

Real data analysis example

- Look at the file inflammation-01.csv
- The data: inflammation level in patients after a treatment (CSV) format
 - Each row holds information for a single patient
 - The columns represent successive days.
 - The first few lines in the file:



Days																																															
Patients	0,0,1,3,1,2,4,7,8,3,3,3,10,5,7,4,7,7,12,18,6,13,11,11,7,7,4,6,8,8,4,4,5,7,3,4,2,3,0,0																																														
	0,1,2,1,2,1,3,2,2,6,10,11,5,9,4,4,7,16,8,6,18,4,12,5,12,7,11,5,11,3,3,5,4,4,5,5,1,1,0,1																																														
	0,1,1,3,3,2,6,2,5,9,5,7,4,5,4,15,5,11,9,10,19,14,12,17,7,12,11,7,4,2,10,5,4,2,2,3,2,2,1,1																																														
	0,0,2,0,4,2,2,1,6,7,10,7,9,13,8,8,15,10,10,7,17,4,4,7,6,15,6,4,9,11,3,5,6,3,3,4,2,3,2,1																																														
	0,1,1,3,3,1,3,5,2,4,4,7,6,5,3,10,8,10,6,17,9,14,9,7,13,9,12,6,7,7,9,6,3,2,2,4,2,0,1,1																																														

Reading the data using NumPy

```
import numpy as np
fname = "inflammation-01.csv"
data = np.loadtxt(fname, delimiter=',')
print(data)
[[ 0.  0.  1. ...,  3.  0.  0.]
 [ 0.  1.  2. ...,  1.  0.  1.]
 [ 0.  1.  1. ...,  2.  1.  1.]
 ...,
 [ 0.  1.  1. ...,  1.  1.  1.]
 [ 0.  0.  0. ...,  0.  2.  0.]
 [ 0.  0.  1. ...,  1.  1.  0.]]
```

properties of the data

```
print(type(data))
print(data.shape)
print("first value in data", data[0,0])
<type 'numpy.ndarray'>
(60L, 40L)
('first value in data', 0.0)
```

Tasks

- Remove the first and last 10 days
- Plot the average, min, and max inflammation score per day


Data trimming and plots

remove the 10 first and last days

n,m = data.shape

data = data[:,10:(m-10)]

How can we get the average/min/max of each day?



Inflammation values per day

```
import matplotlib.pyplot as plt
plt.plot(data.mean(axis=0), label='avg')
plt.plot(data.max(axis=0), label='max')
plt.plot(data.min(axis=0), label='min')
plt.title('inflammation per day')
plt.legend()
plt.show()
```

